

# AI Assistant Architecture Overview


**Date:** January 26, 2025

**Purpose:** System architecture documentation for AI Assistant feature

**Status:** Production Ready

## System Architecture

### High-Level Overview



## Component Structure

### 1. Frontend Component

**File:** components/AIAssistantWidget.tsx

#### Responsibilities:

- Render floating chat widget UI
- Handle user input and message display

- Stream responses from API
- Manage widget state (open/closed/minimized)
- Keyboard navigation (WCAG 2.1 AA)
- Markdown link parsing

#### **Key Features:**

- Client-side component ('use client')
  - Real-time streaming responses
  - Conversation history management
  - Error handling and display
  - Accessibility compliance
- 

## **2. API Route**

**File:** app/api/ai-assistant/route.ts

#### **Responsibilities:**

- Handle POST requests from widget
- Rate limiting (10 requests/minute per IP)
- Validate incoming messages
- Stream responses from Claude API
- Store conversations in Supabase
- Error monitoring with Sentry

#### **Key Features:**

- Streaming responses (Server-Sent Events)
  - Rate limiting with in-memory store
  - Conversation persistence
  - Comprehensive error handling
  - TypeScript type safety
- 

## **3. Claude Integration**

**File:** lib/clause.ts

#### **Responsibilities:**

- Initialize Anthropic SDK client
- Format system prompt from knowledge base
- Stream chat responses
- Handle API errors gracefully
- Validate messages before sending

#### **Key Features:**

- Model: clause-3-haiku-20240307
  - Streaming support
  - Error categorization (402, 403, 429, etc.)
  - Message validation
- 

## **4. Knowledge Base**

**File:** lib/knowledge-base.ts

#### **Responsibilities:**

- Define RWS services and pricing

- Provide FAQ responses
- Format system prompt for Claude
- Include guardrails and guidelines
- Provide website links and CTAs

#### Key Features:

- Centralized pricing (from `lib/config/pricing.ts`)
- Service descriptions
- FAQ database
- Guardrails for off-topic questions
- CTA strategy (questionnaire, booking)

---

## 5. Rate Limiting

File: `lib/rate-limit.ts`

#### Responsibilities:

- Track request counts per IP
- Enforce 10 requests/minute limit
- Return rate limit errors

#### Key Features:

- In-memory store (Map-based)
- IP-based tracking
- Configurable limits

---

## 6. Database Schema

File: `supabase/migrations/20250125_create_ai_assistant_tables.sql`

#### Tables:

- `ai_assistant_conversations` - Conversation metadata
- `ai_assistant_messages` - Individual messages

#### Features:

- UUID primary keys
- Row Level Security (RLS)
- Timestamps (TIMESTAMPTZ)
- Indexes for performance

---

## Data Flow

### 1. User Sends Message

```
User types message
↓
AIAssistantWidget.tsx
↓
POST /api/ai-assistant
↓
Rate limit check
↓
Validate message
```

```
↓  
Store user message in Supabase  
↓  
Call Claude API (streaming)  
↓  
Stream chunks to client  
↓  
Store AI response in Supabase  
↓  
Display in widget
```

## 2. Conversation History

```
Widget loads  
↓  
Check for existing conversationId  
↓  
If exists: Fetch messages from Supabase  
↓  
Display in widget  
↓  
Continue conversation
```

## 🔒 Security & Performance

### Security

1. **Rate Limiting:** 10 requests/minute per IP
2. **Input Validation:** Max 5000 characters, non-empty
3. **RLS Policies:** Supabase Row Level Security enabled
4. **API Key Protection:** Server-side only, never exposed
5. **Error Handling:** No sensitive data in error messages

### Performance

1. **Streaming:** Real-time responses (no waiting for full response)
2. **Lazy Loading:** Widget loads on all pages (consider lazy loading if needed)
3. **Database Indexes:** Optimized queries for conversation history
4. **Error Monitoring:** Sentry for production error tracking

## 🛠️ Technology Stack

### Frontend

- **Next.js 16** (App Router)
- **React 18** (Client components)
- **TypeScript** (Type safety)
- **TailwindCSS** (Styling)

### Backend

- **Next.js API Routes** (Serverless functions)
- **Claude 3 Haiku API** (AI responses)
- **Supabase** (Database & storage)
- **Sentry** (Error monitoring)

### Infrastructure

- **Vercel** (Hosting & deployment)
  - **Supabase Cloud** (Database)
  - **Anthropic Cloud** (AI API)
- 

## Database Schema

### ai\_assistant\_conversations

```
- id (UUID, PRIMARY KEY)
- created_at (TIMESTAMPTZ)
- updated_at (TIMESTAMPTZ)
- user_ip (TEXT)
- metadata (JSONB)
```

### ai\_assistant\_messages

```
- id (UUID, PRIMARY KEY)
- conversation_id (UUID, FOREIGN KEY)
- role (TEXT) -- 'user' or 'assistant'
- content (TEXT)
- created_at (TIMESTAMPTZ)
```

## Configuration

### Environment Variables

```
# Required
ANTHROPIC_API_KEY=sk-ant-...
NEXT_PUBLIC_SUPABASE_URL=https://...
NEXT_PUBLIC_SUPABASE_ANON_KEY=eyJ...
# Optional
SENTRY_DSN=https://... (for error monitoring)
```

### Rate Limiting

**Current:** 10 requests/minute per IP

**Location:** lib/rate-limit.ts

**Storage:** In-memory (resets on server restart)

---

## Deployment

### Vercel Deployment

1. Push to GitHub
2. Vercel auto-deploys
3. Environment variables set in Vercel dashboard
4. Verify widget appears on live site

### Database Migration

1. Run migration: supabase/migrations/20250125\_create\_ai\_assistant\_tables.sql

- 
2. Verify tables created
  3. Check RLS policies enabled
- 

## Monitoring & Analytics

### Sentry

- Error tracking
- Performance monitoring
- User session tracking

### Supabase

- Conversation analytics
- Message counts
- User engagement metrics

### Claude API

- Token usage
  - Response times
  - Cost tracking
- 

## Future Enhancements

### Potential Improvements

1. **Lazy Loading:** Load widget only when user interacts
  2. **React.memo():** Optimize widget re-renders
  3. **Caching:** Cache common responses
  4. **Analytics:** Track conversation quality
  5. **A/B Testing:** Test different system prompts
  6. **Model Upgrade:** Consider Claude 3.5 Sonnet for complex queries
- 

## Related Documentation

- [docs/AI\\_ASSISTANT\\_WIDGET\\_DEPLOYMENT.md](#) - Deployment guide
  - [docs/AI\\_ASSISTANT\\_TESTING\\_GUIDE.md](#) - Testing procedures
  - [docs/MODEL\\_SELECTION\\_GUIDE.md](#) - Haiku vs Sonnet
  - [docs/AI\\_ASSISTANT\\_GUARDRAILS.md](#) - Guardrail implementation
  - [docs/AI\\_ASSISTANT\\_LINKS\\_GUIDE.md](#) - CTA strategy
- 

## Architecture Checklist

- Streaming responses implemented
  - Rate limiting configured
  - Database schema created
  - Error handling comprehensive
  - Accessibility compliant (WCAG 2.1 AA)
  - TypeScript types defined
  - Monitoring configured (Sentry)
  - Documentation complete
- 

**Last Updated:** January 26, 2025

**Status:** Production Ready

**Next Review:** When upgrading model or adding features